

Lab: Filtering a DataFrame

This lab covers:

- Reducing the memory usage of a `DataFrame`
- Extracting a subset of rows from a `DataFrame` based on one or more conditions
- Filtering for rows that include or exclude null values
- Selecting values that fall between a range
- Removing duplicate and null values from a `DataFrame`

Optimizing A Dataset for Memory Usage

Let's begin with the usual import of our favorite data analysis library.

```
1In [1] import pandas as pd
```

The `employees.csv` dataset for this lab is a fictional collection of employees at a company. Each record includes the employee's first name, gender, start date at the firm, salary, management status (True or False), and team. Let's take a peek...

```
12345678910111213141516In [2] pd.read_csv("employees.csv")
```

Out [2]

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	8/6/93	NaN	True	Marketing
1	Thomas	Male	3/31/96	61933.0	True	NaN
2	Maria	Female	NaN	130590.0	False	Finance
3	Jerry	NaN	3/4/05	138705.0	True	Finance
4	Larry	Male	1/24/98	101004.0	True	IT
...
996	Phillip	Male	1/31/84	42392.0	False	Finance
997	Russell	Male	5/20/13	96914.0	False	Product
998	Larry	Male	4/20/13	60500.0	False	Business Dev
999	Albert	Male	5/15/12	129949.0	True	Sales
1000	NaN	NaN	NaN	NaN	NaN	NaN

Is there a way that we can increase the utility of our dataset? Our first optimization is one we should feel pretty comfortable with by now. The text values in the **Start Date** column can be converted to datetime objects with the `parse_dates` parameter.

```
12345678910In [3] pd.read_csv("employees.csv", parse_dates = ["Start Date"]).head()
```

Out [3]

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
3	Jerry	NaN	2005-03-04	138705.0	True	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT

Now that we're in a good place with the import, let's assign the `DataFrame` object to a descriptive variable like `employees`.

```
12In [4] employees = pd.read_csv("employees.csv",
                                parse_dates = ["Start Date"])
```

Converting Data Types with the `astype` Method

There are a few options available for improving the speed and efficiency of operations on the `DataFrame`. The `info` method is particularly helpful here. It returns a big-picture summary of the dataset, including column names, data types, missing values, and memory consumption.

```
12345678910111213In [5] employees.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1001 entries, 0 to 1000
Data columns (total 6 columns):
First Name    933 non-null object
Gender        854 non-null object
Start Date    999 non-null datetime64[ns]
Salary        999 non-null float64
Mgmt          933 non-null object
Team          957 non-null object
dtypes: datetime64[ns](1), float64(1), object(4)
memory usage: 47.0+ KB
```

First up, `Pandas` has imported our **Mgmt** column as a sequence of text values even though it is fundamentally a collection of Booleans. Notice that the designation of the column's values is `object` in the output from the `info` method. `object` is the library's internal lingo for strings.

The `astype` method converts the values of a `Pandas` object to a different data type. Let's extract the **Mgmt** `Series` object from our `DataFrame`, then invoke its `astype` method. The argument to the method represents the data type to convert to. We'll pass `bool`, a built-in keyword in Python. The return value will be a new `Series` object. Note that the `NaN` values have been converted to `True` values.

```
1234567891011121314In [6] employees["Mgmt"].astype(bool)

Out [6] 0      True
        1      True
        2     False
        3      True
        4      True
        ...
       996     False
       997     False
       998     False
       999      True
      1000      True
      Name: Mgmt, Length: 1001, dtype: bool
```

Looks good! It's now time to replace our existing **Mgmt** column.

Let's overwrite our **Mgmt** column with our new `Series` of Booleans. As a reminder, the right side of the assignment operator (=) is always evaluated first.

```
1In [7] employees["Mgmt"] = employees["Mgmt"].astype(bool)
```

Column assignment does not return a value, so we have to output the `DataFrame` again to see the results.

```
12345678910In [8] employees.tail()
```

```
Out [8]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales
1000	NaN	NaN	NaT	NaN	True	NaN

With the exception of the last row of missing values, there are no visual changes to the look and feel of the `DataFrame`. But what about our memory usage? Let's invoke `info` again.

```
1234In [9] employees.info()
```

```
#...
```

```
memory usage: 40.2+ KB
```

Next, let's take a look at the **Salary** column. If we open the raw CSV file, we'll see that the column's values are actually stored as whole numbers. For technical reasons, `Pandas` needs to convert them from integers to floating point values to support the `NaN` values scattered throughout. If we try to coerce the column's values to integers, a `ValueError` exception will be raised.

```
12345678In [10] employees["Salary"].astype(int)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-99-b148c8b8be90> in <module>
----> 1 employees["Salary"].astype(int)

ValueError: Cannot convert non-finite values (NA or inf) to integer
```

Later in the lab, we'll discuss how we can remove this last row entirely. For now, we can use `fillna` method to return a new `Series` where `NaN` values have been replaced with a specified argument. Let's provide an arbitrary value of 0 here. Take a look at the last row.

```
12345678In [11] employees["Salary"].fillna(0).tail()
```

```
Out [11] 996      42392.0
          997      96914.0
          998      60500.0
          999     129949.0
          1000         0.0
          Name: Salary, dtype: float64
```

Now that the **Salary** column has no missing values, we're clear to convert its values to integers...

```
12345678In [12] employees["Salary"].fillna(0).astype(int).head()

Out [12] 0          0
         1      61933
         2     130590
         3     138705
         4     101004
         Name: Salary, dtype: int64
```

...and overwrite the existing `Series` in the `employees DataFrame`.

```
In [13] employees["Salary"] = employees["Salary"].fillna(0).astype(int)
```

The `nunique` method can tell us the number of unique values in each column of a `DataFrame`. Note that missing values (`NaN`) will be excluded by default; you can add a `dropna = False` argument to the `nunique` method to include them in the count.

```
123456789In [14] employees.nunique()

Out [14] First Name      200
         Gender          2
         Start Date     971
         Salary         994
         Mgmt           2
         Team           10
         dtype: int64
```

Two columns stand out as good candidates for categorical values: **Gender** and **Team**. In 1001 rows of data, they only have 2 and 10 unique values respectively.

Let's practice the `astype` syntax again. First, we'll test that a column's values can be converted to a categorical data type.

```
123456789101112131415In [15] employees["Gender"].astype("category")

Out [15] 0          Male
         1          Male
         2        Female
         3          NaN
         4          Male
         ...
        996         Male
        997         Male
        998         Male
        999         Male
       1000         NaN
         Name: Gender, Length: 1001, dtype: category
         Categories (2, object): [Female, Male]
```

Excellent! `Pandas` has identified two unique values, i.e. two categories -- Male and Female. We're good to overwrite our existing **Gender** column and check the memory usage. It has dropped significantly once again because

`pandas` only has to keep track of two values instead of 1001.

```
In [16] employees["Gender"] = employees["Gender"].astype("category")
        employees.info()

#...
memory usage: 33.5+ KB
```

Let's repeat the same process for the **Team** column. There are only ten unique values among 1001 rows.

```
In [17] employees["Team"] = employees["Team"].astype("category")
        employees.info()

#...
memory usage: 27.0+ KB
```

Thanks to these simple methods, we've reduced the memory usage by over 40%! Imagine that impact on datasets with millions of rows!

Filtering by a Single Condition

Let's imagine it's our first day as a data analyst and our manager has asked us to return a list of all company employees named "Maria". To accomplish this, we need to filter our `employees` dataset based on the values in the **First Name** column.

As a reminder, the equal sign (`==`) is used to compare the equality of two objects in Python. For example, we can check if two strings are equal.

```
123In [18] "Maria" == "Maria"

Out [18] True
```

One might think that using the equality operator with a `Series` and a string as operands would lead to an error. After all, they are fundamentally different objects. However, `pandas` is smart enough to recognize that we want to compare *each* value in the `Series` with the specified string. The code below returns a new `Series` of Booleans where a value of `True` indicates a string of "Maria" at that index position in the **First Name** column. For example, Maria is found in the row with index position 2.

```
1234567891011121314In [19] employees["First Name"] == "Maria"

Out [19] 0      False
         1      False
         2       True
         3      False
         4      False
         ...
        996      False
        997      False
        998      False
        999      False
       1000      False
        Name: First Name, Length: 1001, dtype: bool
```

If we can get all the rows with a value of True in the `Series` above selected from our original `employees DataFrame`, we would have all the "Maria" records in the dataset. To accomplish this, pass the Boolean Series in between a pair of square brackets following the `employees` variable.

```
1234567891011In [20] employees[employees["First Name"] == "Maria"]
```

```
Out [20]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
2	Maria	Female	NaT	130590	False	Finance
198	Maria	Female	1990-12-27	36067	True	Product
815	Maria	NaN	1986-01-18	106562	False	HR
844	Maria	NaN	1985-06-19	148857	False	Legal
936	Maria	Female	2003-03-14	96250	False	Business Dev
984	Maria	Female	2011-10-15	43455	False	Engineering

Excellent! If the syntax is slightly confusing due to the use of multiple square brackets, we can also assign the Boolean `Series` to a descriptive variable, then pass *that* variable into the square brackets instead. The code below will yield the same subset of rows as the code above.

```
12In [21] marias = employees["First Name"] == "Maria"
employees[marias]
```

The most common beginner error here is using the wrong number of equal signs. Remember, a single equal sign assigns an object to a variable while two equal signs check for equality between objects.

Let's try another example. What if we wanted to pull out all employees who are *not* on the HR team? We need to generate a Boolean `Series` that checks which values in the **Team** column are *not equal* to "HR".

Python's inequality operator is ideal for comparing that two values are not equal.

```
123In [22] "Engineering" != "HR"
```

```
Out [22] True
```

The `Series` object plays friendly with the inequality operator as well. In the output below, a True indicates that the **Team** value for a given row is *not* "HR", while a False indicates the **Team** value is "HR".

```
In [23] (employees["Team"] != "HR").head()
```

```
Out [23] 0      True
         1     False
         2      True
         3      True
         4      True
         Name: Team, dtype: bool
```

Again, we can pass the `Series` into square brackets to extract the `DataFrame` rows in which an index position has a value of True. Below, we can see that row 1 has been *excluded* because the Team value there is "HR".

```
In [24] employees[employees["Team"] != "HR"]
```

```
Out [24]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
2	Maria	Female	NaT	130590	False	Finance
3	Jerry	NaN	2005-03-04	138705	True	Finance
4	Larry	Male	1998-01-24	101004	True	IT
5	Dennis	Male	1987-04-18	115163	False	Legal
...
995	Henry	NaN	2014-11-23	132483	False	Distribution
996	Phillip	Male	1984-01-31	42392	False	Finance
997	Russell	Male	2013-05-20	96914	False	Product
998	Larry	Male	2013-04-20	60500	False	Business Dev
999	Albert	Male	2012-05-15	129949	True	Sales

866 rows × 6 columns

Note that the results will exclude rows with missing values. A `NaN` value is considered neither equal nor unequal to a string.

What if we wanted to retrieve all of our managers? Do we need to execute `employees["Mgmt"] == True`? We *could* do that but there is no need because we *already* have a `Series` of Booleans. We can just pass the column itself inside the square brackets.

```
In [25] employees[employees["Mgmt"]].head()
```

Out [25]

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
1	Thomas	Male	1996-03-31	61933	True	NaN
3	Jerry	NaN	2005-03-04	138705	True	Finance
4	Larry	Male	1998-01-24	101004	True	IT
6	Ruby	Female	1987-08-17	65476	True	Product

We can also filter columns based on mathematical condition. As long as we provide a Boolean `Series`, `pandas` will be able to filter the `DataFrame`. Let's see which employees earn a six-figure salary.

```
1234567891011In [26] high_earners = employees["Salary"] > 100000
employees[high_earners].head()
```

Out [26]

	First Name	Gender	Start Date	Salary	Mgmt	Team
2	Maria	Female	NaT	130590	False	Finance
3	Jerry	NaN	2005-03-04	138705	True	Finance
4	Larry	Male	1998-01-24	101004	True	IT
5	Dennis	Male	1987-04-18	115163	False	Legal
9	Frances	Female	2002-08-08	139852	True	Business Dev

Filtering by Multiple Conditions

A `DataFrame` can also be filtered by multiple conditions. The strategy is to create two or more Boolean `Series`, then specify the logical criteria that must be met between them.

The AND Condition

Our next ask is to find all female employees working in business development. There are two conditions that must be met for a given row to be selected: a value of "Female" in the **Gender** column and a value of "Business Dev" in the **Team** column. The best strategy here is to construct out one `Series` at a time. We can begin by isolating the "Female" values in the **Gender** column.

```
1In [27] is_female = employees["Gender"] == "Female"
```

Next up, we'll target *all* employees working on the "Business" Dev team.

```
1In [28] in_biz_dev = employees["Team"] == "Business Dev"
```

Finally, we need to calculate the intersection of the two `Series`, the rows in which *both* of them have a value of True. Pass *both* of the `Series` into the square brackets and place an ampersand (`&`) symbol in between them. The `&` specifies an AND criteria. Both `Series` must have a True value at the same index position in order for a row to be selected.

```
12345678910In [29] employees[is_female & in_biz_dev].head()
```

Out [29]

	First Name	Gender	Start Date	Salary	Mgmt	Team
9	Frances	Female	2002-08-08	139852	True	Business Dev
33	Jean	Female	1993-12-18	119082	False	Business Dev
36	Rachel	Female	2009-02-16	142032	False	Business Dev
38	Stephanie	Female	1986-09-13	36844	True	Business Dev
61	Denise	Female	2001-11-06	106862	False	Business Dev

We can pass as many `Series` in the square brackets as we want as long as we separate every subsequent two with a `&` symbol. The next example retrieves all female managers on the "Business Dev" team.

```
1234567891011In [30] is_manager = employees["Mgmt"]
employees[is_female & in_biz_dev & is_manager].head()
```

Out [30]

	First Name	Gender	Start Date	Salary	Mgmt	Team
9	Frances	Female	2002-08-08	139852	True	Business Dev
38	Stephanie	Female	1986-09-13	36844	True	Business Dev
66	Nancy	Female	2012-12-15	125250	True	Business Dev
92	Linda	Female	2000-05-25	119009	True	Business Dev
111	Bonnie	Female	1999-12-17	42153	True	Business Dev

The OR Condition

Rows can also be extracted if they fit one of *several* conditions. For example, what if we wanted to find all employees whose **Salary** is below 40,000 *or* whose **Start Date** was after January 1st, 2015? We can use mathematical operators like `<` and `>` to arrive at two separate Boolean `Series`.


```
12In [31] earning_below_40k = employees["Salary"] < 40000
         started_after_2015 = employees["Start Date"] > "2015-01-01"
```

To specify an OR criteria, use a pipe symbol (`|`). In the next example, a row will be selected if *either* of the Boolean Series at that index position holds a True.

```
12345678910In [32] employees[earning_below_40k | started_after_2015].tail()
```

Out [32]

	First Name	Gender	Start Date	Salary	Mgmt	Team
958	Gloria	Female	1987-10-24	39833	False	Engineering
964	Bruce	Male	1980-05-07	35802	True	Sales
967	Thomas	Male	2016-03-12	105681	False	Engineering
989	Justin	NaN	1991-02-10	38344	False	Legal
1000	NaN	NaN	NaT	0	True	NaN

Inversion with ~

The tilde symbol (`~`) inverts the values in a Series of Booleans. All True values become False, and all False values become True. Here's a simple example.

```
1234567891011121314In [33] my_series = pd.Series([True, False, True])
         my_series
```

Out [33]

0	True
1	False
2	True

dtype: bool

In [34] ~my_series

Out [34]

0	False
1	True
2	False

dtype: bool

This is helpful when we want to invert or reverse a condition. For example, if we wanted to isolate employees earning less than 100,000, we could write `employees["Salary"] < 100000`

```
12345678910In [35] employees[employees["Salary"] < 100000].head()
```

Out [35]

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
1	Thomas	Male	1996-03-31	61933	True	NaN
6	Ruby	Female	1987-08-17	65476	True	Product
7	NaN	Female	2015-07-20	45906	True	Finance
8	Angela	Female	2005-11-22	95570	True	Engineering

...or we could *invert* the results set of employees who earn *more* than (or equal to) 100,000. The results will be identical.

```
12345678910In [36] employees[~(employees["Salary"] >= 100000)].head()

Out [36]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
1	Thomas	Male	1996-03-31	61933	True	NaN
6	Ruby	Female	1987-08-17	65476	True	Product
7	NaN	Female	2015-07-20	45906	True	Finance
8	Angela	Female	2005-11-22	95570	True	Engineering

Methods for Booleans

An alternative syntactical option is available for those who prefer methods over mathematical operators. The table below outlines six built-in methods that all return `Boolean Series`. Note that categorical values do *not* support any mathematical operations besides equality.

+-----+-----+	+-----+-----+	+-----+-----+
Operation	Method Syntax	
+-----+-----+	+-----+-----+	+-----+-----+
Equality	employees["Team"].eq("Marketing")	
+-----+-----+	+-----+-----+	+-----+-----+
Inequality	employees["Team"].ne("Marketing")	
+-----+-----+	+-----+-----+	+-----+-----+
Less than	employees["Salary"].lt(100000)	
+-----+-----+	+-----+-----+	+-----+-----+
Less than or equal to	employees["Salary"].le(100000)	
+-----+-----+	+-----+-----+	+-----+-----+
Greater than	employees["Salary"].gt(100000)	
+-----+-----+	+-----+-----+	+-----+-----+
Greater than or equal to	employees["Salary"].ge(100000)	
+-----+-----+	+-----+-----+	+-----+-----+

Filtering by Condition

Some filtering operations are more complex than a simple equality or inequality check. Luckily, `Pandas` ships with many helper methods that return `Boolean Series`.

The isin Method

What if we wanted to isolate all employees on *either* the Sales, Legal, or Marketing teams? We *could* declare three separate `Series` and use them all inside the square brackets with the OR criteria.

```
In [37] sales = employees["Team"] == "Sales"
        legal = employees["Team"] == "Legal"
        mktg = employees["Team"] == "Marketing"
        employees[sales | legal | mktg].head()

Out [37]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
5	Dennis	Male	1987-04-18	115163	False	Legal
11	Julie	Female	1997-10-26	102508	True	Legal
13	Gary	Male	2008-01-27	109831	False	Sales
20	Lois	NaN	1995-04-22	64714	True	Legal

This approach works but a better solution is the `isin` method, which accepts a list of elements. It returns a Boolean `Series` in which a `True` indicates that a row's value is found amongst the list's values. Once we have the `Series`, we can use it to filter in the usual manner. The next example accomplishes the exact same result as above.

```
123456789101112In [38] all_star_teams = ["Sales", "Legal", "Marketing"]
      in_team = employees["Team"].isin(all_star_teams)
      employees[in_team].head()
```

Out [38]

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
5	Dennis	Male	1987-04-18	115163	False	Legal
11	Julie	Female	1997-10-26	102508	True	Legal
13	Gary	Male	2008-01-27	109831	False	Sales
20	Lois	NaN	1995-04-22	64714	True	Legal

The between Method

Another common challenge, especially when dealing with numeric data, is extracting values that fall within a range. For example, what if we wanted to extract a list of all employees with a salary between 80,000 and 90,000?

Once again, we could use *two* separate `Series`.

```
123456789101112In [39] higher_than_80 = employees["Salary"] >= 80000
      lower_than_90 = employees["Salary"] < 90000
      employees[higher_than_80 & lower_than_90].head()
```

Out [39]

	First Name	Gender	Start Date	Salary	Mgmt	Team
19	Donna	Female	2010-07-22	81014	False	Product
31	Joyce	NaN	2005-02-20	88657	False	Product
35	Theresa	Female	2006-10-10	85182	False	Sales
45	Roger	Male	1980-04-17	88010	True	Sales
54	Sara	Female	2007-08-15	83677	False	Engineering

There's a better way, however. A `Series` object includes a convenient `between` method that accepts a lower bound and upper bound. It returns a Boolean `Series` where a `True` indicates that a row's value falls between the specified interval. Note that the first argument, the lower bound, is inclusive while the second argument, the upper bound, is exclusive. The code below accomplishes the same result as the code above.

```
1234567891011In [40] between_80k_and_90k = employees["Salary"].between(80000, 90000)
      employees[between_80k_and_90k].head()
```

```
Out [40]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
19	Donna	Female	2010-07-22	81014	False	Product
31	Joyce	NaN	2005-02-20	88657	False	Product
35	Theresa	Female	2006-10-10	85182	False	Sales
45	Roger	Male	1980-04-17	88010	True	Sales
54	Sara	Female	2007-08-15	83677	False	Engineering

The `between` method also works on columns of datetime values. We can pass strings representing the start and end dates of our time range. The respective parameters are `left` and `right`. Below, we find all employees who started with the company in the 1980s.

```
123456789101112131415In [41] eighties_folk = employees["Start Date"].between(
                                left = "1980-01-01",
                                right = "1990-01-01"
                                )

                                employees[eighties_folk].head()
```

```
Out [41]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
5	Dennis	Male	1987-04-18	115163	False	Legal
6	Ruby	Female	1987-08-17	65476	True	Product
10	Louise	Female	1980-08-12	63241	True	NaN
12	Brandon	Male	1980-12-01	112807	True	HR
17	Shawn	Male	1986-12-07	111737	False	Product

Finally, we can apply the `between` method to string columns. Let's extract all employees whose first name starts with "R". We'll start with a capital "R" as our inclusive lower bound and go up to the non-inclusive upper bound of "S".

```
1234567891011In [42] name_starts_with_r = employees["First Name"].between("R", "S")
                                employees[name_starts_with_r].head()
```

```
Out [42]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
6	Ruby	Female	1987-08-17	65476	True	Product
36	Rachel	Female	2009-02-16	142032	False	Business Dev
45	Roger	Male	1980-04-17	88010	True	Sales
67	Rachel	Female	1999-08-16	51178	True	Finance
78	Robin	Female	1983-06-04	114797	True	Sales

The isnull and notnull Methods

Our dataset includes plenty of missing values. We can see a few of them in our first five rows.

```
12345678910In [43] employees.head()
```

```
Out [43]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	0	True	Marketing
1	Thomas	Male	1996-03-31	61933	True	NaN
2	Maria	Female	NaT	130590	False	Finance
3	Jerry	NaN	2005-03-04	138705	True	Finance
4	Larry	Male	1998-01-24	101004	True	IT

Missing values are marked with a `NaN` (not a number) designation. The one exception is datetime values, which have a `NaT` (not a time) designation. You can see an example in the **Start Date** column at index position 2.

We can use several methods to isolate rows with either null or present values in a given column. The `isnull` method returns a Boolean `Series` where a `True` indicates that a row's value is absent.

```
In [44] employees["Team"].isnull().head(2)

Out [44] 0    False
         1     True
         2    False
         Name: Team, dtype: bool
```

`NaT` values will be considered null as well.

```
In [45] employees["Start Date"].isnull().head(3)

Out [45] 0    False
         1    False
         2     True
         Name: Start Date, dtype: bool
```

The `notnull` method returns the inverse `Series`, one in which a `True` indicates a row's value is present.

```
In [46] employees["Team"].notnull().head(2)

Out [46] 0     True
         1    False
         Name: Team, dtype: bool
```

This produces the exact same results set as inverting the `Series` returned by the `isnull` method with the tilde (`~`) character. Either approach works, but `notnull` is a bit more descriptive and thus recommended.

```
In [47] (~employees["Team"].isnull()).head(2)

Out [47] 0     True
         1    False
         Name: Team, dtype: bool
```

Once again, we can use these Boolean `Series` to select specific rows from the `DataFrame`.

```
In [48] no_team = employees["Team"].isnull()
         employees[no_team].head()

Out [48]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
1	Thomas	Male	1996-03-31	61933	True	NaN
10	Louise	Female	1980-08-12	63241	True	NaN
23	NaN	Male	2012-06-14	125792	True	NaN
32	NaN	Male	1998-08-21	122340	True	NaN
91	James	NaN	2005-01-26	128771	False	NaN

```
In [49] has_name = employees["First Name"].notnull()
        employees[has_name].tail()
```

Out [49]

	First Name	Gender	Start Date	Salary	Mgmt	Team
995	Henry	NaN	2014-11-23	132483	False	Distribution
996	Phillip	Male	1984-01-31	42392	False	Finance
997	Russell	Male	2013-05-20	96914	False	Product
998	Larry	Male	2013-04-20	60500	False	Business Dev
999	Albert	Male	2012-05-15	129949	True	Sales

Dealing with Null Values

Since we're on the topic of null values, let's discuss some options for dealing with them. Earlier, we saw how we could use the `fillna` method to replace missing values with a constant. Let's bring our dataset back to its original shape by reimporting the CSV. Here's a reminder of what it looks like:

```
In [50] employees = pd.read_csv("employees.csv",
                                parse_dates = ["Start Date"])
```

```
In [51] employees
```

Out [51]

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
3	Jerry	NaN	2005-03-04	138705.0	True	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT
...
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales
1000	NaN	NaN	NaT	NaN	NaN	NaN

1001 rows × 6 columns

By default, the `dropna` method will remove all rows from the `DataFrame` that hold *any* `NaN` values. It doesn't matter if the row has one missing value or six; `dropna` will exclude them all.

```
In [52] employees.dropna()
```

```
Out [52]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
4	Larry	Male	1998-01-24	101004.0	True	IT
5	Dennis	Male	1987-04-18	115163.0	False	Legal
6	Ruby	Female	1987-08-17	65476.0	True	Product
8	Angela	Female	2005-11-22	95570.0	True	Engineering
9	Frances	Female	2002-08-08	139852.0	True	Business Dev
...
994	George	Male	2013-06-21	98874.0	True	Marketing
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales

```
761 rows × 6 columns
```

Improperly exported datasets can often contain blank lines. We can pass an argument of "all" to the `how` parameter of the `dropna` method to remove rows where *all* values are `NaN` or `NaT`. Only one row (our last one at index 1000) satisfies this condition.

```
In [53] employees.dropna(how = "all").tail()
```

```
Out [53]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
995	Henry	NaN	2014-11-23	132483.0	False	Distribution
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales

The `subset` parameter is used to remove rows with a missing value in a specific column. The next example returns a new `DataFrame` consisting only rows that *do not* have a missing value in the **Gender** column.

```
In [54] employees.dropna(subset = ["Gender"]).tail()
```

```
Out [54]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
994	George	Male	2013-06-21	98874.0	True	Marketing
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales

We can also pass the `subset` parameter a list of several strings. A row will be removed if it has a missing value in *any* of the specified columns. The next example returns a new `DataFrame` that removes rows with missing values in *either* the **Start Date** column or the **Salary** column.

```
In [55] employees.dropna(subset = ["Start Date", "Salary"]).head()
```

```
Out [55]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
1	Thomas	Male	1996-03-31	61933.0	True	NaN
3	Jerry	NaN	2005-03-04	138705.0	True	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT
5	Dennis	Male	1987-04-18	115163.0	False	Legal
6	Ruby	Female	1987-08-17	65476.0	True	Product

The `thresh` parameter specifies a minimum threshold of non-null values that a row must have in order to be kept. The next example selects only the rows with at least 4 present values.

```
12345678910In [56] employees.dropna(how = "any", thresh = 4).head()
```

```
Out [56]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
3	Jerry	NaN	2005-03-04	138705.0	True	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT

Dealing with Duplicates

The duplicated Method

In the next example, "Finance" is the value for the **Team** column for the rows at index positions 2 and 3. `Pandas` marks the first occurrence at index 2 as a non-duplicate (with a `False`) and all subsequent occurrences as duplicates (with a `True`).

```
1234567891011121314151617In [57] employees["Team"].head()
```

```
Out [57] 0    Marketing
         1         NaN
         2    Finance
         3    Finance
         4         IT
         Name: Team, dtype: object
```

```
Out [58] employees["Team"].duplicated().head()
```

```
Out [58] 0    False
         1    False
         2    False
         3     True
         4    False
         Name: Team, dtype: bool
```

Conversely, we can ask `Pandas` to mark the *last* occurrence of a value in a column as a non-duplicate. Pass a string of "last" to the `keep` parameter to overwrite its default argument of "first".


```
1234567891011121314In [59] employees["Team"].duplicated(keep = "last")
```

```
Out [59] 0      True
         1      True
         2      True
         3      True
         4      True
         ...
        996    False
        997    False
        998    False
        999    False
       1000    False
Name: Team, Length: 1001, dtype: bool
```

Let's say we wanted to extract one employee from each team. One solution is to pull out the first encountered row for each value in the **Team** column. Our existing `duplicated` method returns a `Series` where a `True` marks all duplicate values after the first encounter. If we *invert* those results, we'll get a `Series` where a `True` marks the *first* time a value is encountered.

```
12345678In [60] (~employees["Team"].duplicated()).head()
```

```
Out [60] 0      True
         1      True
         2      True
         3     False
         4      True
Name: Team, dtype: bool
```

Now we can extract exactly one employee per team. This time around, a `NaN` will be considered a unique value.

```
1234567891011121314151617In [61] first_one_in_team = ~employees["Team"].duplicated()
                                employees[first_one_in_team]
```

```
Out [61]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT
5	Dennis	Male	1987-04-18	115163.0	False	Legal
6	Ruby	Female	1987-08-17	65476.0	True	Product
8	Angela	Female	2005-11-22	95570.0	True	Engineering
9	Frances	Female	2002-08-08	139852.0	True	Business Dev
12	Brandon	Male	1980-12-01	112807.0	True	HR
13	Gary	Male	2008-01-27	109831.0	False	Sales
40	Michael	Male	2008-10-10	99283.0	True	Distribution

The `drop_duplicates` Method

A `DataFrame`'s `drop_duplicates` method provides a convenient shortcut to accomplish the operations above. By default, it will remove any rows where *all* values are shared with a previously encountered row. Because the combination of 6 values in each row is unique in our dataset, it won't accomplish anything on its own.

```
12345678910111213141516In [62] employees.drop_duplicates()

Out [62]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
3	Jerry	NaN	2005-03-04	138705.0	True	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT
...
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales
1000	NaN	NaN	NaT	NaN	NaN	NaN

Much like with the `deduplicated` method, the `subset` parameter can specify a list of columns whose values will be used to determine a row's uniqueness. The next example again finds the *first* occurrence of each value in the **Team** column.

```
12345678910111213141516In [63] employees.drop_duplicates(subset = ["Team"])

Out [63]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT
5	Dennis	Male	1987-04-18	115163.0	False	Legal
6	Ruby	Female	1987-08-17	65476.0	True	Product
8	Angela	Female	2005-11-22	95570.0	True	Engineering
9	Frances	Female	2002-08-08	139852.0	True	Business Dev
12	Brandon	Male	1980-12-01	112807.0	True	HR
13	Gary	Male	2008-01-27	109831.0	False	Sales
40	Michael	Male	2008-10-10	99283.0	True	Distribution

The `drop_duplicates` method also accepts a `keep` parameter. We can pass in an argument of "last" to keep the rows with the *last* occurrence of each encountered value.

```
12345678910111213141516In [64] employees.drop_duplicates(subset = ["Team"], keep =
"last")

Out [64]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
988	Alice	Female	2004-10-05	47638.0	False	HR

989	Justin	NaN	1991-02-10	38344.0	False	Legal
990	Robin	Female	1987-07-24	100765.0	True	IT
993	Tina	Female	1997-05-15	56450.0	True	Engineering
994	George	Male	2013-06-21	98874.0	True	Marketing
995	Henry	NaN	2014-11-23	132483.0	False	Distribution
996	Phillip	Male	1984-01-31	42392.0	False	Finance
997	Russell	Male	2013-05-20	96914.0	False	Product
998	Larry	Male	2013-04-20	60500.0	False	Business Dev
999	Albert	Male	2012-05-15	129949.0	True	Sales
1000	NaN	NaN	NaT	NaN	NaN	NaN

The `keep` parameter accepts one other argument, `False`, which will discard *all* rows that have duplicate values. The next example select all rows in the dataset where the value in the **First Name** column is unique.

```
1234567891011121314In [65] employees.drop_duplicates(subset = ["First Name"], keep =
False)

Out [65]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
5	Dennis	Male	1987-04-18	115163.0	False	Legal
8	Angela	Female	2005-11-22	95570.0	True	Engineering
33	Jean	Female	1993-12-18	119082.0	False	Business Dev
190	Carol	Female	1996-03-19	57783.0	False	Finance
291	Tammy	Female	1984-11-11	132839.0	True	IT
495	Eugene	Male	1984-05-24	81077.0	False	Sales
688	Brian	Male	2007-04-07	93901.0	True	Legal
832	Keith	Male	2003-02-12	120672.0	False	Legal
887	David	Male	2009-12-05	92242.0	False	Legal

If the list passed to the `subset` parameter includes two or more strings, `Pandas` will consider the combo of those values across a row when identifying a duplicate. For example, here is a list of all employees with a **First Name** of "Douglas" and a **Gender** of "Male".

```
1234567891011In [66] name_is_douglas = employees["First Name"] == "Douglas"
is_male = employees["Gender"] == "Male"
employees[name_is_douglas & is_male]

Out [66]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
217	Douglas	Male	1999-09-03	83341.0	True	IT
322	Douglas	Male	2002-01-08	41428.0	False	Product
835	Douglas	Male	2007-08-04	132175.0	False	Engineering

If we pass a list of `["Gender", "Team"]` to the `subset` parameter of the `drop_duplicates` method, the combination of values across those two columns will be used to determine duplicates. The row at index 0 represents the first occurrence of "Douglas" as a **First Name** value and "Male" as a **Gender** value. If a row had a **First Name** value of "Douglas" and a **Gender** value that was *not* "Male", it would be included. Similarly, if had a row had a **Gender** value of "Male" but a **First Name** value that was *not equal* to "Douglas", it would be included as well. It is the

combination of the two that `Pandas` uses to identify duplicates. Thus, the first row from the output above will be included while the others will be excluded.

```
12345678910In [67] employees.drop_duplicates(subset = ["Gender", "Team"]).head()
```

```
Out [67]
```

	First Name	Gender	Start Date	Salary	Mgmt	Team
0	Douglas	Male	1993-08-06	NaN	True	Marketing
1	Thomas	Male	1996-03-31	61933.0	True	NaN
2	Maria	Female	NaT	130590.0	False	Finance
3	Jerry	NaN	2005-03-04	138705.0	True	Finance
4	Larry	Male	1998-01-24	101004.0	True	IT

Coding Challenge

The Problem

We've tackled some heavy data work for our company's HR department. Now's your chance to tackle something a bit less corporate. The `netflix.csv` file is a collection of almost 6,000 titles available to watch in November 2019 on the popular online video streaming service Netflix. The **director** and **date_added** columns contain missing values (we can see examples in the rows at index positions 0, 2, and 5836 below).

```
12345678910111213141516In [68] pd.read_csv("netflix.csv")
```

```
Out [68]
```

	title	director	date_added	type
0	Alias Grace	NaN	3-Nov-17	TV Show
1	A Patch of Fog	Michael Lennox	15-Apr-17	Movie
2	Lunatics	NaN	19-Apr-19	TV Show
3	Uriyadi 2	Vijay Kumar	2-Aug-19	Movie
4	Shrek the Musical	Jason Moore	29-Dec-13	Movie
...
5832	The Pursuit	John Papola	7-Aug-19	Movie
5833	Hurricane Bianca	Matt Kugelman	1-Jan-17	Movie
5834	Amar's Hands	Khaled Youssef	26-Apr-19	Movie
5835	Bill Nye: Science Guy	Jason Sussberg	25-Apr-18	Movie
5836	Age of Glory	NaN	NaN	TV Show

Using the skills in this lab, answer the following questions about the dataset.

1. How we can we optimize the dataset for speed and utility?
2. Find all rows with a title of "Limitless"
3. Find all rows with a director of "Robert Rodriguez" and a type of "Movie"
4. Find all rows with *either* a date_added of "2019-07-31" or a director of "Robert Altman"
5. Find all rows with a director of "Orson Welles", "Aditya Kripalani" or "Sam Raimi".
6. Find all rows that have a date_added value between May 1st, 2019 and June 1st, 2019
7. Drop all rows with a NaN value in the director column.
8. Identify all days when only *one* movie was added to the Netflix catalog

5.6.2 Solutions

Let's tackle the questions!

1. How we can we optimize the dataset for speed and utility?

For utility, it's optimal to store the values in the **date_added** column as datetime objects. We can force this type coercion on CSV import.

```
1In [69] netflix = pd.read_csv("netflix.csv", parse_dates = ["date_added"])
```

Let's take a look at the current memory usage.

```
123456In [70] netflix.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5837 entries, 0 to 5836
#...
memory usage: 182.5+ KB
```

Can any column's values be converted to categorical values? Let's use the `nunique` method to count the number of unique values per column.

```
1234567In [71] netflix.nunique()

Out [71] title          5780
         director      3024
         date_added    1092
         type           2
         dtype: int64
```

The **type** column is about as perfect a candidate as possible. In a dataset of 5837 rows, there are only two unique values in **type**: "Movie" and "TV Show". Let's convert its values with the `astype` method.

```
1In [72] netflix["type"] = netflix["type"].astype("category")
```

How much this has reduced our memory usage? A whopping 22%!

```
1234In [73] netflix.info()

#...
memory usage: 142.7+ KB
```

2. Find all rows with a title of "Limitless"

To solve this one, we need to compare each value in the **title** column to the string of "Limitless". Then, we can select the rows for which that evaluation returns True.

```
12345678In [74] netflix[netflix["title"] == "Limitless"]

Out [74]

         title          director date_added  type
1559  Limitless      Neil Burger  2019-05-16  Movie
2564  Limitless                NaN  2016-07-01  TV Show
4579  Limitless  Vrinda Samarth  2019-10-01  Movie
```

3. Find all rows with a director of "Robert Rodriguez" and a type of "Movie"

This problem requires two `Series`, one comparing the values in the **title** column to the string of "Robert Rodriguez" and the other comparing the values in the **type** column to the string "Movie". The `&` symbol can be used to apply AND logic.

```
12345678910111213In [75] directed_by_rr = netflix["director"] == "Robert Rodriguez"
      is_movie = netflix["type"] == "Movie"
      netflix[directed_by_rr & is_movie]
```

Out [75]

	title	director	date_added	type
1384	Spy Kids: All the Time in the ...	Robert Rodriguez	2019-02-19	Movie
1416	Spy Kids 3: Game...	Robert Rodriguez	2019-04-01	Movie
1460	Spy Kids 2: The Island of Lost D...	Robert Rodriguez	2019-03-08	Movie
2890	Sin City	Robert Rodriguez	2019-10-01	Movie
3836	Shorts	Robert Rodriguez	2019-07-01	Movie
3883	Spy Kids	Robert Rodriguez	2019-04-01	Movie

4. Find all rows with *either* a date_added of "2019-07-31" or a director of "Robert Altman"

This problem is similar to the previous one but requires a `|` symbol for OR logic.

```
123456789101112In [76] added_on_july_31 = netflix["date_added"] == "2019-07-31"
      directed_by_altman = netflix["director"] == "Robert Altman"
      netflix[added_on_july_31 | directed_by_altman]
```

Out [76]

	title	director	date_added	type
611	Popeye	Robert Altman	2019-11-24	Movie
1028	The Red Sea Diving Resort	Gideon Raff	2019-07-31	Movie
1092	Gosford Park	Robert Altman	2019-11-01	Movie
3473	Bangkok Love Stories: Innocence	NaN	2019-07-31	TV Show
5117	Ramen Shop	Eric Khoo	2019-07-31	Movie

5. Find all rows with a director of "Orson Welles", "Aditya Kripalani" or "Sam Raimi".

One option here is to create three separate Boolean `Series`, each one comparing the values of the **director** column with one of the three strings. But the better strategy is to use the `isin` method on the **director** column and pass in a list with the values. It's more concise and scalable.

```
12345678910111213In [77] directors = ["Orson Welles", "Aditya Kripalani", "Sam
      Raimi"]
      target_directors = netflix["director"].isin(directors)
      netflix[target_directors]
```

Out [77]

	title	director	date_added	type
946	The Stranger	Orson Welles	2018-07-19	Movie
1870	The Gift	Sam Raimi	2019-11-20	Movie
3706	Spider-Man 3	Sam Raimi	2019-11-01	Movie

4243	Tikli and Laxmi Bomb	Aditya Kripalani	2018-08-01	Movie
4475	The Other Side of the Wind	Orson Welles	2018-11-02	Movie
5115	Tottaa Pataaka Item Maal	Aditya Kripalani	2019-06-25	Movie

6. Find all rows that have a `date_added` value between May 1st, 2019 and June 1st, 2019

The most concise way to solve this problem is to use the `between` method with the two dates as the lower and upper bounds. This saves us the need to declare two separate `Series`.

```
12345678910111213In [78] may_movies = netflix["date_added"].between(
    "2019-05-01", "2019-06-01"
)
netflix[may_movies].head()
```

Out [78]

	title	director	date_added	type
29	Chopsticks	Sachin Yardi	2019-05-31	Movie
60	Away From Home	NaN	2019-05-08	TV Show
82	III Smoking Barrels	Sanjib Dey	2019-06-01	Movie
108	Jailbirds	NaN	2019-05-10	TV Show
124	Pegasus	Han Han	2019-05-31	Movie

7. Drop all rows with a NaN value in the director column.

The `dropna` method can remove any rows in the `DataFrame` with a NaN value. We just have to pass its `subset` parameter the columns to look for null values in.

```
12345678910In [79] netflix.dropna(subset = ["director"]).head()
```

Out [79]

	title	director	date_added	type
1	A Patch of Fog	Michael Lennox	2017-04-15	Movie
3	Uriyadi 2	Vijay Kumar	2019-08-02	Movie
4	Shrek the Musical	Jason Moore	2013-12-29	Movie
5	Schubert In Love	Lars Büchel	2018-03-01	Movie
6	We Have Always Lived in the Castle	Stacie Passon	2019-09-14	Movie

8. Identify all days when only *one* movie was added to Netflix

There are a couple ways to solve this problem. One solution is to recognize that the **`date_added`** column will have duplicate date values for any titles that were added on the same day. We can invoke the `drop_duplicates` method with the `keep` parameter set to `False` to remove any rows that contain duplicates in the **`date_added`** column.

```
12345678910111213141516In [80] netflix.drop_duplicates(subset = ["date_added"], keep
= False)
```

Out [80]

	title	director	date_added	type
4	Shrek the Musical	Jason Moore	2013-12-29	Movie
12	Without Gorky	Cosima Spender	2017-05-31	Movie

30	Anjelah Johnson: Not Fancy	Jay Karas	2015-10-02	Movie
38	One Last Thing	Tim Rouhana	2019-08-25	Movie
70	Marvel's Iron Man & Hulk: Heroes ...	Leo Riley	2014-02-16	Movie
...
5748	Menorca	John Barnard	2017-08-27	Movie
5749	Green Room	Jeremy Saulnier	2018-11-12	Movie
5788	Chris Brown: Welcome to My Life	Andrew Sandler	2017-10-07	Movie
5789	A Very Murray Christmas	Sofia Coppola	2015-12-04	Movie
5812	Little Singham in London	Prakash Satam	2019-04-22	Movie